

Multiple Phases with Due Dates in Canvas: 1) Proposal in Canvas forum ASAP; 2) Draft-stage completion in GitHub & Presentations in class or pre-recorded video; 3) Peer feedback and voting on projects; 4) All final-revision files in GitHub for grading.

Overview and High-Level Requirements:

Many previous students have done nice work, of which they and I were proud. In a few cases, their projects and this course directly helped them get a job. You have an even wider range of choices available than most earlier semesters, so avoid the pitfalls, apply yourselves the best you can and try to have some fun with it.

1. You must choose to *work as a team of 2 or 3 students*. All teams must clearly show collaboration from every member and larger teams are expected to perform at a proportionally-increased level of complexity, sophistication, depth, and/or scope to earn a high grade. Projects will be evaluated and graded as a teamwork product. You must work together, and hold each other accountable for contributions, quality, and ethical behavior. If there's a problem early on (such as someone insists on plagiarizing, notify the instructor immediately).
2. Select from one of these TYPES of analytics projects to implement in Python:
 - I. **This should be the type that most of you choose!** A critique and improvement of a published data analysis. Includes identification of weaknesses, mistakes, or scope limitations in the published analysis. Your code rewrite/enhancement of the program to improve the rigor of the analysis AND make the code more reliable & maintainable. This is like Assignment 6 except you'll do more work to expand on the original.
 - II. Your original Monte Carlo simulation.
 - III. Your original analysis linking 2+ published data sets from distinct sources, investigating some topic like: societal or environmental changes possibly affected by changes in laws, industry, new inventions, or corporate practices; or analysis of biases within data; or testing scientific hypotheses; or uncovering evidence of corruption in any industry, company, or government; historical changes of health, economic, or other aggregate life quality factors between countries, cities, or similar.
3. Do not propose any statistical Machine Learning-focused project. Enroll in the Machine Learning Team Projects course if that interests you. Thus, packages such as sklearn, PyTorch, or TensorFlow and similar should not be used here. NLP-focused projects are also out of scope – take Text Mining instead.
4. You must submit your original unique work, created specifically for 597PR. If the project is related to work you did earlier or are now doing for any other course or a job, then you must get prior written approval from all the relevant instructors and the job supervisor. Not doing so is subject to sanctions per the Student Code.
5. PROPOSAL Stage. Post your summary into Canvas's **"Final Project Proposals"** forum. See expected information there.
6. CITE ALL YOUR SOURCES! Any citation style is fine, but make sure you do it and do so continuously, not just at the very end. Students who use code or other material without indicating its true source AND clearly delineating which parts are not original will be reported to the iSchool & UI Grad

College through FAIR. If the review confirms plagiarism, the sanction is reduced grade (dependent on scope) with a course grade of F for significant cases, regardless of your previous scores.

7. Unlike all other assignments in 597PR, you are allowed and expected to openly publish your unique project work. You may consider it part of your student portfolio, link it from resume, etc. Typically, this is done on Github.com since you'll be committing there as work proceeds. Make sure everything you put there is work you'll be proud of. At every commit, you should be verifying citations for all code, data, etc. Remember every commit is a "version" of the project and is public.
8. Draft Stage & Peer Reviews: The program and documentation should be sufficiently operational for meaningful and beneficial peer code review, but does not have to be 100% finished. Make sure your GitHub repository is up to date with all the work you've done so far (code, documentation, example outputs). It's okay if there are some final scenario explorations or even minor flaws left to resolve in your project at this stage, but you want constructive feedback from others. All students will also be submitting evaluations about other teams' projects, details will be given in class.
9. Presentations: You will create and deliver a presentation to the class (either live in person or as pre-recorded video in Illinois MediaSpace) that summarizes your project's purpose, hypotheses, design reasoning, and results so far, possibly with a quick demo.
10. Submission expectations:
 - ☐ Edit your README.md to create a good introduction and overview of the project, written with new visitors to your repository in mind. Summarize the conclusions you came up with, including how results are either supportive of or refute your hypotheses. [You can embed images](#) into the README file, if that is relevant.
 - ☐ Use the "Factors in Code Quality and Code Reviews" like a checklist! Apply as many of the skills we've discussed this semester as applicable, to create the best quality program you can. Example expectations:
 - i. Doctests and/or pytest unit tests: 2-person teams minimum 50% actual test coverage; (3-person minimum 80% actual coverage).
 - ii. All functions need complete Docstrings and use type annotations properly. 2-person=7+ functions; 3-person= 10+ functions
 - iii. 3-person teams also should incorporate one of these efficiency techniques that will be discussed in class later: Selective compilation (e.g. Numba or Cython); and/or parallel processing.
 - ☐ Consider each hypothesis or alternative situation you proposed to investigate (you may have added more after feedback). Your program should be able to automatically run the experiments and outputs for ALL of the hypotheses -- do not hard-code such configuration aspects into the functions themselves in a way that requires hand-editing to use or adjust the program.

(Type I Projects) Specifics for a Formal Critique and Improvement of a Published Data Analysis:

The original publication you critique and rework does not have to be a scholarly peer-reviewed analysis but it can be – or a part of one. So, it could be a less formal analysis that was published on an open website (e.g. a blog, GitHub.com, Kaggle.com, etc.) or in a magazine (like your Assignment 6 was but you expand on it more).

I don't want to see a project where you essentially just convert an already-good analysis from another language into Python. The original might even be in Python already, but it needs to be flawed, fragile, or incomplete in its code, statistical analysis, and/or conclusions. That could have happened because they used biased or insufficient data (that you will improve, augment, or work around), made logical mistakes during analysis, or based conclusions on incorrect programming. Another more subtle possibility is that they improperly ignored uncertainties within the raw or processed data.

Most of the notebooks ("Code") on Kaggle are low-quality code and many have weak or even pointless analysis, both of which leave a lot of room for you to improve them. **But** the popular topics and data sets on Kaggle can have hundreds of posts with code by different people using the same data. That density doesn't leave much room for improvement without plagiarism. So, whether it's on Kaggle or elsewhere, you must choose a data analysis topic to improve that does NOT already have dozens or hundreds of people's versions and commentary on it.

(Type II projects) Specifics for Monte Carlo Simulations:

- ☐ Design your own scenario -- make certain your simulation is original in some way(s).
- ☐ It MUST clearly show all 3 Phases of a complete MC Sim as defined in my video.
- ☐ You can simulate an engineering or manufacturing problem, business/management situation, (certain types of) human behaviors, physical phenomena, or a game. To encourage original thinking, **AVOID scenarios that have been done many times and/or discussed in class, such as** : a "random walk" of stock prices, stock options, or similarly naïve financial "predictions"; a traffic simulation with just a few intersections or only one road; parking lots or parking meters; customer seating/dining at a restaurant or serving them at a counter; the games *Tic-tac-toe*, *four*-in-a-row*, *Go-moku*, *chutes and ladders*, *Monopoly*, *Rock-Paper-Scissors*, *Blackjack*, *Poker*. If you want to model a sports game or tournament ask first, as too many of these have been done already.
- ☐ Phase 1 - You must have several well-chosen random variables in the model, to explore a variety of possible outcomes and derive the non-obvious probabilities of the overall model. Make sure you think carefully to choose appropriate ranges of values and a sensible distribution type for every randomized aspect. For example, if you simulate "number of swimmers in the pool" at each point in time as a uniform distribution, it's wrong. If you simulate the individual finish times of all runners in a marathon as uniform, triangular, or even normal, it's wrong. We'll discuss this in class.
 - i. 2-person teams must have at least 3 different kinds of randomized variables plus the deterministic aspects and control variables in the model. Most interesting simulations require more.
 - ii. 3-person teams must have at least 4 different kinds ...
- ☐ If there is any relevant public data available, try to incorporate real data as part of your simulation model. If data you seek is not in downloadable form, you can still research the scenario to make realistic estimates.
- ☐ Phase 2 & 3 -- Controls & Experiments: You need to demonstrate running the simulation in a "control" situation to help establish that it works properly before doing the experiments. There should be multiple hypotheses to test by running your code with options to create different.

- Phase 3 – Analysis. Explain the aggregate results you've generated, whether and how it supports or refutes your hypotheses for each experiment. If Phase 1 or 2 weren't done well, document and revise as needed to improve and run new variation(s).

(Type III Projects) Specifics for an Original Data Analysis [Non-simulation]:

An **original** analysis linking 2+ published data sets from DIFFERENT sources, investigating some topic like: societal or environmental changes over time (possibly) affected by changes in laws, industry, new inventions, or corporate practices; or analysis of complex biases within data; or testing scientific hypotheses; or showing evidence of corruption in any industry, company, or government; historical changes of health, economic, or other aggregate life quality factors between countries, cities, or similar.

There are *thousands* of public data sets that could be of interest to you, from many US and foreign government agencies, scientific organizations, universities, companies, and more. Look at data.gov, worldbank.org, or similar big repositories where you can search for open data that interests you. But to do original work using popular data sets typically requires *combining* them meaningfully with very different data sources. Type III proposals most often get rejected because they don't meet this criterion.

Some earlier homework assignments were like this in concept, but your final project should be larger in scope, sophistication, and code quality than the regular assignments you did in 597PR.

Warning: many previous plagiarism cases in this course have come from students copying chunks of analysis and program code from one or more Kaggle.com or similar data analysis found on GitHub and presenting it as their own work. If you do the same, from ANY uncredited source, you will likely fail this course. If something on Kaggle strongly interests you, consider a Type I project on it instead.